

Exercise metric learning : find your doppleganger

Module 5, Master Computer Vision, 2018-19

Joan Serrat

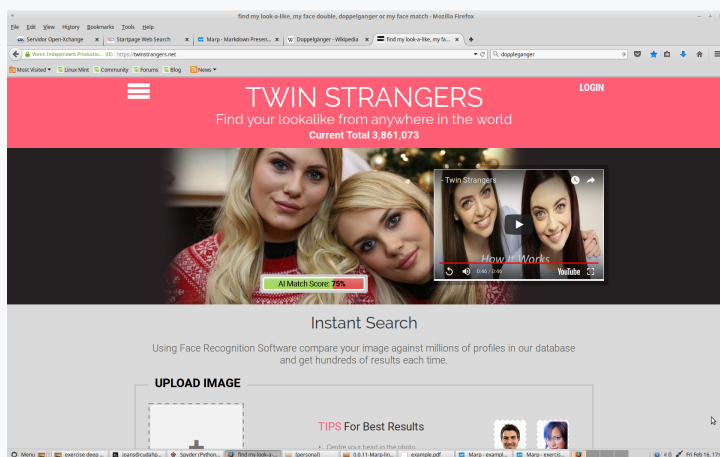
<http://www.cvc.uab.es/people/joans/>

What's a doppelganger ?

A person who closely resembles another but is unrelated to him/her. Also, a double and a look-alike. Can be a living person, a portrait in a museum ...



Iranian doppelganger of Leo Messi



<https://twinstrangers.net/>

How to find him/her ?

- define a **metric**, ie. a measure of similarity
- or better, **learn it** from samples
- samples = **pairs of images** from one the same person
- dataset with **many persons, many images** per identity
- once learned, provide an image and find the most similar instance in the dataset

This is impossible with the available data and time. So we'll just do our best on a small dataset

It won't be a real look-alike, but hopefully will somehow resemble you!

The *real* objectives of the exercise

- Learn **by example** how to :

- build a siamese/triplet network
- draw pairs of samples to train it
- make a simple distance-based classifier

in **TensorFlow**, **Keras** or **PyTorch**, as you prefer

- Implement a **loss** function
- Evaluate the effect of some **mining** technique on the training and accuracy
- **Visualize** the embedding with t-SNE
- Perform and assess **image retrieval**

Datasets : Labeled Faces in the Wild


- +5K identities of celebrities, sportsmen, politicians ... (back in 2009)
- highly variable number of images per identity
- only 62 identities with > 20 images and 168 with > 10 images
- **cropped face region** to avoid influence of background
- color, **resized** 64 x 64

Datasets : Labeled Faces in the Wild



62 identities

Datasets : Labeled Faces in the Wild



168 identities

Datasets : Tsinghua Traffic Signs

- +170 traffic signs from China
- diverse camera orientation, changing lighting, occlusions
- cropped from the original frames and resized 64 x 64
- only 81 traffic signs with > 20 images

Datasets : Tsinghua Traffic Signs



81 traffic signs

What to do

1. Learn a **baseline** metric for Tsinghua and LFW for a dimension of 2 (just for visualization) and then 16, 32 or 64
2. Compute classification **accuracy** in Tsinghua with some distance-based classifier : nearest class mean, k -nearest neighbors (implemented in scikit-learn.org)
3. **Visualize** the training and testing set samples on the 2d and n d embedding spaces in Tsinghua and comment the results. Use PCA or t-SNE in TensorBoard or scikit-learn.org
4. **Do retrieval** : find your k doppleganger candidates in LFW as k -NN ($k = 1$ is ok)
5. Change the contrastive loss to **triplet loss**
6. Do *some kind* of **mining** and compare with baseline in terms of accuracy and evolutions of the loss. Tensorflow has some already implemented.

What to do

This is not a challenge to achieve the best accuracy.

I've got 95% in Tsinghua (81 classes) and just 83% (62 classes) in LFW

It's more about how to

- **learn an embedding and visualize it**
- **do mining**
- **use it for classification and explain the results**

Code

We provide a lot of code ready to use :

- siamese based on VGG-16
- generation of batches of pairs for both Tsinghua and LFW
- split datasets into train / validation / test
- accuracy and nearest neighbor computation
- visualization in 2d and n d with scikit-learn t-SNE (but TensorBoard implementation is much cooler)
- one tower + siamese loss, hint on triplet loss

Code + datasets [here](#)

Grading

B = 0.25 exam points if you do the basics : 1 to 4

A = 0.5 points if B + 5 or B + 6

Deliverables

Source code and at most 10 slides in PDF packed in a zip file to be sent to joans@cvc.uab.es by April 2nd, 23:59.

This is team work. Tell me who did what in the last slide.